

2.

$$\begin{aligned} E[V(x)] &= E[x - \frac{\alpha x^2}{2}] = E[x] - \frac{\alpha}{2} E[x^2] \\ &= M - \frac{\alpha}{2} (\text{Var}(x) + E[x]^2) = M - \frac{\alpha}{2} (b^2 + M^2) \end{aligned}$$

$$X_{CE} \text{ s.t. } V(x_{CE}) = E[V(x)] = M - \frac{\alpha}{2} (b^2 + M^2)$$

$$X_{CE} - \frac{\alpha X_{CE}^2}{2} = M - \frac{\alpha}{2} (b^2 + M^2)$$

$$-\frac{\alpha}{2} X_{CE}^2 + X_{CE} - (M - \frac{\alpha}{2} (b^2 + M^2)) = 0$$

$$X_{CE} = 1 \pm \sqrt{1 - 2\alpha(M - \frac{\alpha}{2} (b^2 + M^2))}$$

$$\bar{\pi}_A = E[x] - X_{CE} = M - \frac{1 \pm \sqrt{1 - 2\alpha(M - \frac{\alpha}{2} (b^2 + M^2))}}{\alpha}$$

$$\approx -\frac{1}{2} \cdot \frac{V''(M)}{V'(M)} \cdot b^2 = \frac{1}{2} \frac{\alpha}{1-\alpha M} b^2$$

$$\frac{-V''(x)}{V'(x)} = \frac{\alpha}{1-\alpha x}$$

$$A_A(x) = \frac{\alpha}{1-\alpha x} \quad X_{CE} \approx \bar{x} - \bar{\pi}_A = M - \frac{1}{2} \frac{\alpha}{1-\alpha M} b^2$$

$$\pi = \frac{2}{10^6} \quad W \sim N(1+r + \pi(M-r), \pi^2 b^2)$$

$$\text{maximize } E[V(W)] = M_W - \frac{\alpha}{2} (b_w^2 + M_w^2)$$

$$\therefore 1+r + \pi(M-r) - \frac{\alpha}{2} (\pi^2 b^2 + (1+r + \pi(M-r))^2) \approx 1+r + \pi(M-r)$$

$$\frac{\partial E[V(W)]}{\partial \pi} = M_r - b^2 \alpha \pi - \alpha(1+r + \pi(M-r))(M-r) = 0$$

$$M_r - b^2 \alpha \pi - \alpha(1+r)(M-r) - \alpha \pi (M-r)^2 = 0$$

$$M_r - \alpha(1+r)(M-r) = b^2 \alpha \pi - \alpha \pi (M-r)^2$$

$$M_r - \alpha(1+r)(M-r) = \pi(\alpha(b^2 - (M-r)^2))$$

$$\bar{\pi}^* = \frac{M_r - \alpha(1+r)(M-r)}{\alpha(b^2 - (M-r)^2)} \Rightarrow z^* = 10^6 \bar{\pi}^*$$

$$\text{with } M=0.2, r=0.05, b^2=10,000$$

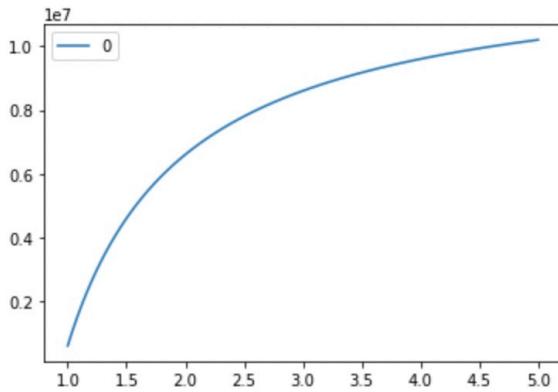
```
In [43]: import numpy as np
import pandas as pd

In [51]: mu = 0.2
r = 0.05
sigma2 = 0.01
z = [1e6 * (mu - r - a * (1 + r) * (mu - r)) / (a * (sigma2 - (mu - r) ** 2)) for alpha in np.linspace(1, 5, 100)]

In [47]: plot_df = pd.DataFrame(z_list, index=alpha)

In [48]: plot_df.plot()

Out[48]: <AxesSubplot:>
```



3. - $W_1 = W_0(1 + f(1+\alpha))$
 - $W_1 = W_0(1 + f(1-\beta))$
 • - $V(W_1) = \log(W_0) + \log(1 + f(1+\alpha))$
 - $V(W_1) = \log(W_0) + \log(1 + f(1-\beta))$
 • $E[\log(W_1)] = p\log(W_0) + \log(1 + f(1+\alpha)) + (1-p)\log(W_0) + \log(1 + f(1-\beta))$
 $= \log(W_0) + p\log(1 + f(1+\alpha)) + (1-p)\log(1 + f(1-\beta))$
 • $\frac{\partial E[\log(W_1)]}{\partial f} = \frac{p(1+\alpha)}{1+f(1+\alpha)} + \frac{(1-p)(1-\beta)}{1+f(1-\beta)}$
 • $\frac{\partial E[\log(W_1)]}{\partial f} = 0 \Rightarrow \frac{p(1+\alpha)}{1+f(1+\alpha)} = \frac{(1-p)(1-\beta)}{1+f(1-\beta)}$
 $(1+f(1-\beta))p(1+\alpha) = -(1+f(1+\alpha))(1-p)(1-\beta)$
 $p(1+\alpha) + (1-p)(1-\beta) = -f(1+\alpha)(1-p)(1-\beta) - f(1-\beta)p(1+\alpha)$
 $p(1+\alpha) + (1-p)(1-\beta) = -f(1+\alpha)(1-p)(1-\beta) + (1+\alpha)p(1-\beta)$
 $f^* = \frac{-p(1+\alpha) - (1-p)(1-\beta)}{(1+\alpha)(1-\beta)}$
 $f^* = \frac{-p}{1-\beta} - \frac{(1-p)}{1+\alpha}$

$$\frac{\partial^2 E[\log(W_1)]}{\partial f^2} = \frac{-p(1+\alpha)^2}{(1+f(1+\alpha))^2} - \frac{(1-p)(1-\beta)^2}{(1+f(1-\beta))^2}$$

evaluate for $f = f^*$:

$$\frac{-p(1+\alpha)^2}{\left(1 + \left(\frac{-p}{1-\beta} - \frac{(1-p)}{1+\alpha}\right)(1+\alpha)\right)^2} - \frac{(1-\alpha)(1-\beta)^2}{\left(1 + \left(\frac{-p}{1-\beta} - \frac{(1-p)}{1+\alpha}\right)(1-\beta)\right)^2}$$

$$\frac{(-)}{(+)}) - \frac{(+)}{(+)}) = (-), \text{ so } f^* \text{ indeed a maxima}$$

- $f^* \uparrow$ as $\beta \downarrow$; increase wager when loss potential decreases
- $f^* \uparrow$ as $\alpha \uparrow$; increase wager when gain potential increases
- $f^* \uparrow$ as $p \uparrow$ ($\text{if } \beta > 1$): increase wager as probability for gain increases

1.

```
In [112]: from typing import Iterator, Tuple, TypeVar, Sequence, List, Dict
from operator import itemgetter
import numpy as np
import os
import sys
sys.path.append(os.path.abspath("/Users/justincramer/Documents/Coding/CME241/RL-book/"))
import itertools
from rl.distribution import Distribution, Choose
from rl.function_approx import FunctionApprox, Tabular
from rl.iterate import iterate
from rl.markov_process import (FiniteMarkovRewardProcess, MarkovRewardProcess,
                                RewardTransition, NonTerminal, State)
from rl.markov_decision_process import (FiniteMarkovDecisionProcess,
                                         MarkovDecisionProcess,
                                         StateActionMapping)
from rl.policy import DeterministicPolicy, FinitePolicy, FiniteDeterministicPolicy
import rl.approximate_dynamic_programming
from rl.approximate_dynamic_programming import evaluate_mrp, value_iteration, extended_vf, evaluate_finite_mrp
from rl.dynamic_programming import policy_iteration as exact_policy_iteration
```

```
In [37]: S = TypeVar('S')
A = TypeVar('A')

# A representation of a value function for a finite MDP with states of
# type S
ValueFunctionApprox = FunctionApprox[NonTerminal[S]]
QValueFunctionApprox = FunctionApprox[Tuple[NonTerminal[S], A]]
NTStateDistribution = Distribution[NonTerminal[S]]
```

```
In [98]: # approximate policy iteration
def policy_iteration(
    mdp: MarkovDecisionProcess[S, A],
    γ: float,
    approx_0: Tuple[ValueFunctionApprox[S], FinitePolicy[S, A]],
    non_terminal_states_distribution: NTStateDistribution[S],
    num_state_samples: int
) -> Iterator[Tuple[ValueFunctionApprox[S], FinitePolicy[S, A]]]:
    def update(vf_policy: Tuple[ValueFunctionApprox[S], FinitePolicy[S, A]])\
            -> Tuple[ValueFunctionApprox[S], FiniteDeterministicPolicy[S, A]]:
        vf, pi = vf_policy
        # apply policy to get mrp
        mrp: FiniteMarkovRewardProcess[S] = mdp.apply_finite_policy(pi)
        vf_iter: Iterator[ValueFunctionApprox[S]] = evaluate_mrp(mrp, γ, vf,
                                                               non_terminal_states_distribution,
                                                               num_state_samples)
        vf: ValueFunctionApprox[S] = list(itertools.islice(vf_iter, 200))[-1]

        nt_states: Sequence[NonTerminal[S]] = \
            non_terminal_states_distribution.sample_n(num_state_samples)

        def return_(s_r: Tuple[State[S], float]) -> float:
            s1, r = s_r
            return r + γ * extended_vf(vf, s1)

        d: Dict[S, A] = {s: max(mdp.actions(s),
                               key=lambda a: mdp.step(s, a).expectation(return_)) for s in nt_states}

        pi: FiniteDeterministicPolicy[S, A] = FiniteDeterministicPolicy(d)

        return vf, pi

    v_0, pi_0 = approx_0
    return iterate(update, (v_0, pi_0))
```